


Sigmoid Function



$$P(\text{win}) = \sigma(\mathbf{w} \cdot \mathbf{x} + b)$$

where $\sigma(z) = \frac{1}{1 + e^{-z}}$, $\mathbf{x}$ = feature vector, $\mathbf{w}$ = learned weights

- **Implemented from scratch in Python with Numpy**
gradient descent optimization, L2 regularization
- **Leave-one-season-out cross-validation**
17 training seasons, 1 test season, repeated for 18 seasons
- **1,000 epochs per run**